

The Next Repertoire

Designing Your Professional Skillset for AI-Augmented Practice

WORKING PREMISE The unit of design is not the AI tool. It is the relation among professional behavior, task, data, technology, organization, and consequences.

David J. Cox, Ph.D., M.S.B., BCBA-D | Ryan L. O'Donnell, M.S., BCBA
3-hour applied workshop | Up to 3 Learning CEUs, subject to ACE Provider confirmation

PROJECT CHIRON

A practical workbook for selecting, governing, and testing one AI-augmented workflow.



01 • ORIENT

How to use this workbook

PRIVACY BOUNDARY Use entirely fictional examples. Do not enter client, caregiver, employee, payer, or other identifying information into workshop tools. Optional polls use broad categories; facilitators do not retain individual responses.

By the end of the workshop, you will be able to ...

- 1. Differentiate descriptive, predictive, and generative AI by what each system receives, does, produces, and leaves for a professional to decide.
- 2. Analyze an AI-augmented ABA workflow using an act-in-context canvas that makes stakeholders, data boundaries, scope, decision rights, and feedback loops explicit.
- 3. Evaluate a likely use case against privacy and security, validation and generalizability, scope, monitoring and drift, training, vendor, payer, and build-or-buy considerations.
- 4. Produce a 30-day prototype that combines at least one technical safeguard with one organizational behavior management (OBM) contingency.

CE structure

Time	Instructional block	What you produce
0:00–0:50	1 • See the system	A professional repertoire hypothesis
0:50–1:00	Break 1	Return to your assigned interest table
1:00–1:50	2 • Use cases, governance, Define, Ideate	A defined use case and alternative options
1:50–2:00	Break 2	—
2:00–3:00	3 • Ideate, prototype, Field Lab, close	A tested workflow and practice plan

01A • DESIGN THINKING

The process we will practice

Design thinking is an iterative, human-centered approach to understanding a problem, generating alternatives, and learning from a tangible test. We use the familiar five modes; governance runs through all of them.

Five modes

Empathize: learn about people’s work, needs, and constraints through observation and conversation.

Define: synthesize a focused point of view and an open “How might we” question.

Ideate: generate alternatives before choosing against explicit criteria.

Prototype: make a small version of the workflow that people can inspect.

Test: observe what happens, revise, and revisit earlier modes when needed.

What this workshop can and cannot establish

Peer interviews illuminate professional experience. They do not substitute for client, caregiver, or other stakeholder input. Label assumptions and identify whose perspective is missing. A table exercise can expose weaknesses in a design; it cannot establish clinical effectiveness or authorize deployment.

Our route

00:00–00:27 Introduction and AI overview

00:27–00:42 Capture, map, and compare professional repertoires

00:42–00:50 Use cases, examples, and table assignments

00:50–01:00 Break 1; return to your assigned interest table

01:00–01:27 Use-case exercise and governance; 01:27–01:41 Define

01:41–01:50 Ideate; 01:50–02:00 Break 2; 02:00–02:25 Ideate and prototype

02:25–02:50 Chiron Field Lab; 02:50–03:00 commit and close

Follow the page number and section code shown on each writing prompt. Pages follow the teaching sequence; section codes stay stable across the printed activity materials. Later activities may refer back to completed pages. “Everyone writes” means your own workbook; “Table recorder” means one shared answer.

Source: Stanford d.school, Design thinking bootleg (see section 17).

02 • NOTICE

AI by function

For each system, ask: What goes in? What transformation occurs? What comes out? Who acts next?

Function	Plain-language job	ABA-adjacent example	Professional obligation
Descriptive	Organize or summarize what has already happened	Graph patterns; summarize notes	Check completeness, context, and representation
Predictive	Estimate a future value or category	Forecast attrition or outcome risk	Check local validity, uncertainty, scope, and consequences
Generative	Produce new text, images, code, or options	Draft goals, summaries, training examples	Verify facts, logic, fit, authorship, and confidential data handling
Workflow wrapper	Connect a model to forms, records, or actions	Route approvals; prefill documentation	Map the whole process, access controls, failure paths, and human override

WATCH THE CATEGORY ERROR A fluent output is not evidence that a recommendation is accurate, appropriate, ethical, or effective in this context.

My current working definition of artificial intelligence

Write your current understanding in one or two sentences. You'll have an opportunity to revise it after the overview.

My revised working definition of artificial intelligence

Write your current understanding in one or two sentences. You'll have an opportunity to revise it after the overview.

One example I want to test today

02A • AI LITERACY

Terms and trade-offs worth keeping

Terms

Model: a system with learned statistical relationships. A large language model (LLM) processes and generates language sequences. A token is a unit of input or output, sometimes only part of a word.

Prompt and context: the instructions and information supplied for a request. The context window limits the information available to that request.

Training changes learned parameters. Inference uses a trained model. Input reuse, logging, retention, and further training depend on the service and contract.

Retrieval supplies selected external material to a request. Fine-tuning changes parameters through further training. Neither guarantees accuracy.

Generated content can contain unsupported claims, omissions, or fabricated citations. Fluency is not evidence of authority. Verify claims against appropriate sources and context.

Using, configuring, or building

Buying a product, configuring a workflow around a model, fine-tuning, and training from scratch are different undertakings. Compare ongoing review, security, integration, training, monitoring, and exit costs. Open weights or local hosting do not automatically establish privacy or fitness. Include a non-AI option.

Environmental and economic effects

Training and inference both use resources. Data centers also support non-AI workloads. Energy and water demands depend on the workload, hardware, location, and infrastructure; avoid a universal per-prompt claim. Ask whether the task warrants the system, its scale, and the frequency of calls (IEA, 2025).

Protecting professional skill

De-skilling means an established repertoire weakens; *never-skilling* means a needed repertoire is not acquired. These are possible effects of arrangements that remove meaningful practice and feedback. They are not inevitable outcomes of AI. Preserve opportunities to perform, observe, and receive feedback.

Trade-offs to examine

- Speed < > verification
- Access < > privacy
- Scale < > local fit
- Automation < > skill practice.

Choose one of the four trade-offs above, for example speed versus verification. Which of its two sides does your work currently favor, and what would have to change for the other side to win?

03A • EMPATHIZE

Capture experience and compare conditions

Your output: one reported work need, one assumption, and one missing perspective.

Optional role/setting snapshot (1 minute). Keep it broad and anonymous.

Partner interview (4 minutes). Switch who is speaking after 2 minutes.

Individual capture (2 minutes). Separate what was reported from what you inferred.

Use entirely fictional examples. Do not enter identifying or rare combinations of personal details.

Interview prompts

Describe a recurring professional task. Where does work get stuck? What does effective performance look like? Who experiences the consequences? Ask about work conditions rather than collecting case details.

One observed need or recurring work pattern

One assumption and the stakeholder perspective still missing

Compare after the quadrant exercise (complete page 7 before completing)

The section 03 quadrants use the same dimensions as “good at/not yet good at” and “like/dislike doing.” Treat self-ratings as hypotheses. They are not competency assessments. Compare one difference with a partner. What changes across setting, role, workload, or feedback?

To strengthen _____, I need _____
in the context of _____.

Possible roles overlap: Designer, Implementer, and Supporter. Disliking a task does not make it appropriate to automate; liking it does not mean the current arrangement works well.

03 • MAP

Your professional repertoire

A role is a pattern of behavior supported by particular contexts, skills, and consequences, and it can change as those change.

Role	What the repertoire includes	Useful question
Designer	Build, configure, train, evaluate, or connect models and data pipelines	Can I make the system fit the task and its constraints?
Implementer	Integrate a system into clinical or operational work	Can I redesign the workflow and preserve accountable judgment?
Supporter	Define problems, contribute data, test, report failures, and shape requirements	Can I make the use case and feedback useful to builders?

Governance is a responsibility across all three roles. It cannot be handed off to a separate role.

Repertoire map

	ENERGIZING	DRAINING
STRONG NOW	Behaviors I can contribute immediately	Behaviors I can perform, but may not sustain
DEVELOPING	Skills worth deliberately practicing	Consider supports, practice, or redesign

04 • CHOOSE

Select a problem worth working on

Choose the table where the outcome matters and the workflow is familiar enough to analyze.

Table	Likely use case	Question to define
A	Clinical documentation	How can drafting reduce latency while keeping source verification and authorship explicit?
B	Behavior-plan drafting	How can a qualified reviewer notice missing evidence without delegating clinical authorization?
C	Goal recommendations	How can options support a collaborative choice without replacing assessment or preferences?
D	Supervision and training	How can practice and feedback improve observed performance rather than only verbal fluency?
E	Outcome prediction	What evidence would make a prediction useful for this task, population, setting, and decision?
F	Scheduling and workflow	How can scheduling improve efficiency while preserving access, continuity, and meaningful override?

TABLE ROUND Introduce yourself with role, setting, and a fictional workshop handle. Name one measurable work problem. Products come later.

My table and problem statement

Table recorder: In [context], AI may [bounded action] to support [purpose]. [Human role] retains responsibility for [decision].

- ☐ The outcome is socially significant.
- ☐ The workflow occurs often enough to practice.
- ☐ A small pilot is feasible within 30 days.
- ☐ We can identify who benefits and who carries risk.

07 • GOVERN

A lifecycle for responsible adoption

Governance changes as the system moves from selection to deployment, monitoring, and retirement.

Lifecycle stage	Required behaviors	Evidence to retain
1 • Select	Define the problem; compare options; examine evidence, scope, data practices, vendor terms, payer and regulatory context	Requirements; decision record; risk review
2 • Deploy	Set policy and decision rights; disclose appropriately; train; pilot; establish access, escalation, and rollback	Approved workflow; training; pilot criteria
3 • Monitor + audit	Measure outcomes and errors; sample outputs; evaluate impact; watch version changes and drift; receive incident reports	Performance log; incidents; audits; changes
4 • Change + retire	Revalidate after material change; pause when thresholds fail; export or destroy data; communicate replacement or exit	Stop decision; data disposition; transition

CASP PRACTICE PARAMETER AI should support, rather than supplant, clinical work. Practitioners remain responsible for professional decisions, informed implementation, training, monitoring, and impact (Council of Autism Service Providers, 2025).

Our minimum viable governance

- ☐ Named owner
- ☐ Data boundary
- ☐ Pilot measure
- ☐ Stop threshold
- ☐ Approved use and prohibited use
- ☐ Human verification
- ☐ Error-report path
- ☐ Retirement/data plan

07A • GOVERNANCE CHECK

Governance Check: Five Questions Across Eight Domains

Use this during the governance block and keep it as a reference. A missing answer can be a reason to delay use. Name the person or team who can supply evidence or authority.

Privacy and security

Map source, access, retention, reuse, disclosure, deletion, and downstream services. Section 08 holds the data-flow worksheet.

Model validation and scope

Specify task, population, setting, data, version, and consequence. Compare with a meaningful baseline. Generated claims and predictive estimates require different checks. Section 09 records evidence and limits.

Model and data drift

Watch input distributions, missingness, workflow changes, version changes, error, and burden. Name an owner, cadence, and condition for re-evaluation, restriction, or retirement. An error is not automatically evidence of drift.

Operational training and general AI literacy

Identify the behavior required. Instruct, model, rehearse, observe, and provide feedback. Check performance under normal workload. Policy acknowledgments and quiz scores are weak evidence on their own. Sections 10–11 connect authority and behavior.

Vendor conversations and build or buy

Ask for task-specific evidence, data terms, change notice, incident response, export/deletion, and exit support. Compare total maintenance burden of buying, configuring, fine-tuning, training, or using no AI. Section 14 contains the procurement checklist.

One unresolved question, the owner who can answer it, and the decision it blocks

Source: NIST AI RMF and Generative AI Profile (section 17).

08 • SAFEGUARD

Privacy, security, and confidentiality

A de-identification claim is only as good as the process, context, and residual re-identification risk.

WORKSHOP RULE No real protected health information, personally identifiable information, workforce-sensitive information, or payer credentials in workshop tools.

Ask before data enter	Why it matters
Is the tool approved for this information and purpose?	Consumer access is not the same as an organizationally approved, contractually governed service.
Does a service provider create, receive, maintain, or transmit electronic protected health information (ePHI)?	HIPAA obligations may apply even when a provider cannot view encrypted data; a compliant business associate agreement (BAA) and safeguards may be required.
Who can access, reuse, train on, disclose, or retain the data?	Settings, contracts, integrations, logs, and subprocessors may matter more than the user interface.
Can the minimum necessary information accomplish the task?	Reduce exposure; prefer fictional, synthetic, or properly de-identified data for practice and testing.
How will access end and data be exported or destroyed?	Exit is part of the design from the start.

Sketch the data path

Source	What enters	Where processed/stored	Who receives output	Retention/exit

This worksheet is educational. It is not legal advice. Apply organizational policy and obtain qualified privacy, security, compliance, or legal review when needed.

09 • VERIFY

Validity, scope, drift, and error

Ask “For which task, people, setting, data, version, threshold, and consequence does it work?” rather than “Does it work?”

Before pilot	During pilot	After deployment
Define target behavior/outcome and comparison	Sample outputs against a human-reviewed standard	Monitor outcomes, errors, complaints, overrides, and version changes
Specify intended and prohibited use	Test boundary cases and likely failure modes	Revalidate when population, workflow, data, model, or threshold changes
Identify affected groups and context	Inspect differences that could create inequitable impact	Pause, constrain, retrain, replace, or retire when thresholds fail

Define a minimum test

Measure	Current baseline	Pilot criterion	Sample/interval	Owner
Benefit/outcome				
Error/quality				
Equity/access				
Time/workload				

STOP RULE If _____ exceeds/falls below _____ for _____, then _____ will pause or constrain the use case while _____ reviews it.

10 • ASSIGN

Human decision rights

Automation is not one switch. Separate generating, recommending, approving, acting, documenting, and monitoring.

Task step	AI may ...	Human must ...	Evidence / record
Generate options			
Summarize evidence			
Recommend or rank			
Approve decision			
Implement action			
Evaluate outcome			

ACCOUNTABILITY TEST If the output is wrong, who detects it, before what consequence, using what evidence, with what authority to stop the workflow?

Statements we can defend

The AI system is permitted to ...

The AI system is not permitted to ...

14 • PROCURE

Questions for a vendor, or your own build team

A polished demonstration is not a substitute for evidence, contracts, operational fit, and an exit plan.

- ☐ What exact problem and intended use does the system support?
- ☐ What data were used to develop and evaluate it?
- ☐ For which populations, settings, languages, and tasks is performance known?
- ☐ What error measures, uncertainty, limitations, and boundary conditions are documented?
- ☐ What client or workforce data enter the system, including logs and metadata?
- ☐ Where are data processed and stored; who are the subprocessors?
- ☐ Can data or outputs be reused to train models or improve services? Can that be disabled?
- ☐ What agreement, including a BAA when applicable, governs the service?
- ☐ How are access, encryption, incident response, deletion, and audit supported?
- ☐ How are model, prompt, policy, and product changes communicated and versioned?
- ☐ Can we export our records and terminate use without workflow lock-in?
- ☐ What monitoring, human override, escalation, rollback, and retirement functions exist?

BUILD, BUY, OR DO NEITHER Compare total implementation burden (i.e., validation, training, oversight, workflow redesign, integration, monitoring, incident response, and exit) rather than subscription price alone.

05 • DEFINE

Use-Case Canvas: purpose and context

Describe the ongoing act in context before you describe the technology.

1

Aim. What socially significant outcome should improve? What is the smallest useful unit of change?

2

Act + context. Whose behavior is involved? In which workflow, setting, constraints, and contingencies?

3

Stakeholders. Who benefits, decides, uses, receives, oversees, pays, or may be adversely affected?

4

Current performance. What happens now? What evidence shows the problem, its frequency, cost, latency, quality, or risk?

Use-Case Canvas: model and data

Make the data path and the model's job observable enough to inspect.

5

Model function. What enters? What transformation occurs? What output appears? What is the uncertainty or error mode?

6

Data boundary. What information enters; where is it stored; who can access, use, disclose, or retain it; and how is it deleted?

7

Evidence + scope. For whom, where, and for what task has performance been demonstrated? Where should use stop?

8

Decision rights. What may the system suggest, draft, rank, or automate? Who verifies, approves, documents, owns, and can override?

05A • DEFINE PRACTICE

Define a need and revise it

Your output: a person-need-insight statement and an open “How might we...?” question for your table’s case.

Use the shared example below as a model and apply its structure to your table’s case. Draft, get feedback, then revise when a condition changes.

14 minutes total. Use the canvas on pages 15–16 as a reference.

Worked example (Shared demonstration only; apply the structure to your table’s case)

A fictional supervision team observes that new staff receive little timely feedback on performance. The need is timely, observable feedback. An “AI trainer” is a product, and products come after the need is clear. How might we increase useful feedback without hiding errors or treating lesson completion as competence? Mark which observations are established and which require inquiry.

First pass and Rep 1

[Person in context] needs [meaningful need] because [evidence-grounded insight].

How might we...? Keep more than one possible solution available.

Rep 2

A partner names one ordinary overlooked condition. Revise the need or boundary and explain what changed.

Changed condition: _____

Revised definition: _____

Missing voice or evidence: _____

Feedback check: Is the person/context clear? Is the need meaningful? Is the insight supported or labeled as an assumption? Does the question avoid selecting a tool in advance?

11 • DESIGN

Pair a technical solution with an OBM solution

A control that exists only in software, or only in policy, will likely fail when work becomes busy.

Design layer	Examples	Our choice
Technical	Access controls; data minimization; templates; required fields; confidence display; sampled review; version lock; logging; rollback	
Behavioral (OBM)	Task clarification; skills practice; prompts; modeling; checklists; performance feedback; reinforcement; escalation; workload redesign	
Measurement	Accuracy; completeness; latency; rework; override; incident; acceptability; client outcome; equity; opportunity cost	

Behavioral systems sketch

Antecedents / context	Target performer behavior	Consequences / feedback	Measure

IMPLEMENTATION QUESTION What contingencies make the safe behavior easier, faster, more likely, and more reinforcing than the unsafe shortcut?

11A • IDEATION PRACTICE

Generate and compare alternatives

Your output: distinct options and the staff supports each would need.

Generic example: To improve handoffs, compare an AI summary, a checklist, and a revised handoff routine. Each needs a reviewer and a check for missing information.

Before the break: generate options below and get feedback.

After the 10-minute break: design staff supports below. Continue on page 21 to consider consequences and select an option; prepare the prototype on page 22.

Rep 1 Technical and non-AI alternatives

Write silently for 1 minute, combine ideas for 3, then group distinct options for 1. Keep at least one non-AI option. Defer judgment while generating.

Option A and the mechanism it changes

Option B and the mechanism it changes

Non-AI alternative

Rep 2 OBM arrangements

Name the performer, observable response, opportunity, and consequence. Pair each candidate with the time, training, rehearsal, feedback, authority, and workload support it needs. “Human review” alone is not a design.

When _____, [performer] will _____.

The supporting consequence/feedback will be _____.

Feedback received and one revision

12 • ANTICIPATE

Scan for second-order effects

The first-order effect is what the tool does. The second-order effect is how behavior and systems change because the tool is present (e.g., who checks less, who waits longer, which skills go unpracticed).

Domain	Possible benefit	Possible cost or shifted burden	What we will observe
Client choice + autonomy			
Clinical quality			
Staff skill / deskilling			
Workload + pace			
Access + equity			
Documentation + payer			
Environmental resources			

DESKILLING CHECK Which professional discriminations, judgments, or practice opportunities could weaken if this task is routinely delegated, and how will we preserve them?

11A • IDEATION PRACTICE CONTINUED

Trace effects and select a test

Rep 3 Second-order effects

Trace a possible downstream consequence through at least two links. Use section 12 to consider client choice, quality, skill, workload, access, documentation, and resources. Treat these as hypotheses to examine.

Immediate change → next consequence → effect on people or work

Who gains, who bears new work or risk, and who loses practice?

What will we observe, and what revision follows?

Converge using explicit criteria

Compare at least two candidates for purpose/stakeholder fit, evidence, burden, human authority, and a small reversible test. Popularity or novelty alone is not a reason. A non-AI or not-yet decision is legitimate.

Selected option and reason

Assumption most likely to change our choice

Prepare the prototype

In section 13, walk one ordinary fictional case through input, AI contribution, review, decision, and release/monitoring. Specify benefit and error/burden measures plus a stop condition. The Field Lab continues this prototype; it does not introduce a new case.

13 • PROTOTYPE

Build the smallest testable workflow

A prototype should be cheap enough to change and clear enough to fail informatively.

Prototype element	Decision
User and setting	
Trigger / start condition	
Minimum data input	
AI function and output	
Required human verification	
Action or handoff	
Outcome and error measures	
Escalation and stop rule	

Test record and next test

Rep	Input case	Predicted result	Observed result	Change before next rep
1				
2				
3				

13A • CHIRON FIELD LAB

Run the 25-minute prototype and test

Keep your A–F use case and prototype. There are no points or winners. The purpose is to expose assumptions and improve the relation among task, people, technology, and consequences.

Roles and materials

Navigator reads and paces; Recorder moves notes and records revisions; Challenger asks for evidence. Roles may be shared. Use the System Map board, sticky notes, Operations card, matching sealed envelope, and eight decision tokens.

3 AI MAY = permitted contributions. 3 HUMAN MUST = required judgment or action. 2 STOP / ESCALATE = detectable condition, responsible person, and next step. Counts scaffold this exercise; they are not a rule for clinical design. Multiple tokens can share a zone. An unused/prohibited AI role may be explicitly marked.

Sequence

00–02 Launch: Keep your use case. Assign Navigator, Recorder, and Challenger.

02–06 Map: Move the workflow you already designed onto the board using sticky notes.

06–10 Place: Use 3 AI MAY + 3 HUMAN MUST + 2 STOP / ESCALATE tokens. Name one safeguard, benefit measure, and error/burden measure.

10–11 Reveal: Open only your matching sealed context-shift envelope. Read it aloud.

11–18 Revise: Move at least one token. Revise a safeguard, measure, or stop rule. Explain why.

18–21 Challenge: Exchange a 60-second report with a neighboring table. Each asks one question.

21–25 Debrief: Two selected tables report for 60 seconds each. Presenters synthesize the revisions.

Keep envelopes sealed until the presenter cues the reveal. Foil marks the reveal boundary; it is not a score or approval badge. Use sticky notes on the board rather than wet-erase markers.

The table tent communicates WORKING, NEED A PROMPT, or READY TO REPORT. Continue with your best bounded assumption while waiting; presenter approval is not required to proceed.

13A • FIELD LAB RECORD

Capture the revision and its reason

Table family A–F: _____ Purpose: _____

Before the reveal: AI role, human decision, and technical/OBM safeguard

What changed in the envelope, and which assumption no longer holds?

After the reveal: token moved, design revision, and why it matters

Benefit measure: _____ Error/burden measure: _____

Stop condition, responsible person, and next step

Five stems for the 60-second report

Our purpose is... _____

AI may... _____

A human must... _____

We will measure... _____

We stop or escalate when... _____

Neighbor's question or unresolved evidence need

A useful test can lead to revise, constrain, choose a non-AI option, or not deploy. This fictional tabletop exercise does not validate a clinical intervention.

15 • COMMIT

Your 30-day repertoire experiment

Choose a small professional behavior you can practice with feedback. A vague intention to “use more AI” will not work.

Element	My commitment
Use case and socially significant aim	
My likely role: Designer / Implementer / Supporter	
Skill or discrimination to practice	
Three practice opportunities	1. 2. 3.
Person or source providing feedback	
One technical safeguard	
One OBM contingency	
Benefit and error measures	
Stop / reconsider threshold	
Review date	

COMMITMENT STATEMENT By _____, I will _____ in the context of _____, measure _____, and pause or revise if _____.

16 • RETAIN

Field guide: eight questions before use

Keep this page. Use it whenever an AI-enabled workflow moves from curiosity to practice.

1. Purpose: What socially significant outcome are we trying to improve, and for whom?
2. Context: What ongoing act, workflow, constraints, contingencies, and stakeholders matter?
3. Function: What information enters, what does the model do, and what output appears?
4. Data: Is the tool approved for these data, who can use or retain them, and how do we exit?
5. Scope: For whom and where has performance been demonstrated; where should use stop?
6. Decision rights: What may the system draft or recommend, and what must a qualified person verify and own?
7. Implementation: What technical safeguards, skills, prompts, feedback, workload, and escalation support safe use?
8. Monitoring: Which benefits, errors, inequities, version changes, and stop thresholds will we track?

PRAGMATIC CRITERION Does this arrangement support effective action toward an explicit purpose in this context, and can it detect and correct error before meaningful harm?

Knowledge check

1. Name one difference between generative and predictive AI that changes how you would validate an output.

2. Identify one professional decision that should remain explicitly assigned to a qualified human in your use case.

3. State one benefit measure, one error measure, and one stop rule for your prototype.

17 • REFERENCES

Selected references and practice anchors

- Behavior Analyst Certification Board. (2020). Ethics code for behavior analysts. <https://www.bacb.com/ethics-information/ethics-codes/>
- Behavior Analyst Certification Board. (2026). Authorized Continuing Education provider handbook. <https://www.bacb.com/continuing-education-providers/>
- Council of Autism Service Providers. (2025). Practice parameters for the use of artificial intelligence in applied behavior analysis. https://assets-002.noviams.com/novi-file-uploads/casp/pdfs-and-documents/AI_Guidelines/aipracticeparameters.pdf
- Cox, D. J., & Jennings, A. M. (2024). The promises and possibilities of artificial intelligence in the delivery of behavior analytic services. *Behavior Analysis in Practice*, 17, 123–136. <https://doi.org/10.1007/s40617-023-00864-3>
- Hayes, S. C., Hayes, L. J., & Reese, H. W. (1988). Finding the philosophical core: A review of Stephen C. Pepper's World Hypotheses. *Journal of the Experimental Analysis of Behavior*, 50(1), 97–111. <https://doi.org/10.1901/jeab.1988.50-97>
- Jennings, A. M., & Cox, D. J. (2024). Starting the conversation around the ethical use of artificial intelligence in applied behavior analysis. *Behavior Analysis in Practice*, 17, 107–122. <https://doi.org/10.1007/s40617-023-00868-z>
- National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0) (NIST AI 100-1). <https://doi.org/10.6028/NIST.AI.100-1>
- National Institute of Standards and Technology. (2024). Artificial intelligence risk management framework: Generative artificial intelligence profile (NIST AI 600-1). <https://doi.org/10.6028/NIST.AI.600-1>
- U.S. Department of Health and Human Services. (n.d.). Guidance on HIPAA & cloud computing. <https://www.hhs.gov/hipaa/for-professionals/special-topics/health-information-technology/cloud-computing/index.html>
- Stanford d.school. (n.d.). Design thinking bootleg. <https://dschool.stanford.edu/tools/design-thinking-bootleg>
- International Energy Agency. (2025). Energy and AI. <https://www.iea.org/reports/energy-and-ai>

SCOPE NOTE This workbook is a reusable educational tool. It does not substitute for clinical judgment, organizational policy, payer requirements, professional standards, privacy and security review, or legal advice.